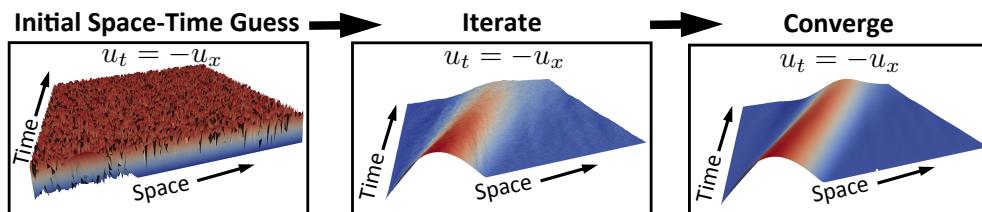


Warp

Parallel-in-Time Software



V. A. Dobrev, R. D. Falgout, Tz. V. Kolev, N. A. Petersson, J. B. Schroder, U. M. Yang
Center for Applied Scientific Computing (CASC)
Lawrence Livermore National Laboratory

This work performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344. LLNL*****

Copyright (c) 2013, Lawrence Livermore National Security, LLC. Produced at the Lawrence Livermore National Laboratory. Written by the WARP team. All rights reserved.

This file is part of WARP. Please see the COPYRIGHT and LICENSE file for the copyright notice, disclaimer, contact information and the GNU Lesser General Public License.

WARP is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License (as published by the Free Software Foundation) version 2.1 dated February 1999.

WARP is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the IMPLIED WARRANTY OF MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the terms and conditions of the GNU General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

Contents

1 Abstract	1
2 Introduction	1
2.1 Introduction	1
2.2 Overview of Warp Algorithm	1
2.2.1 Two-Grid Algorithm	5
2.2.2 Summary	5
2.3 Overview of Warp Code	6
2.3.1 Parallel decomposition and memory	6
2.3.2 Cycling and relaxation strategies	7
2.3.3 Overlapping communication and computation	8
2.3.4 Configuring the Warp Hierarchy	8
2.3.5 Heat equation example	9
2.4 Summary of Warp	10
3 Example	10
4 Building Warp	15
5 Compiling and running the examples	15
6 Module Index	16
6.1 Modules	16
7 File Index	16
7.1 File List	16
8 Module Documentation	16
8.1 User-written routines	16
8.1.1 Detailed Description	17
8.1.2 Typedef Documentation	17
8.2 User interface routines	19
8.2.1 Detailed Description	19
8.2.2 Typedef Documentation	19
8.2.3 Function Documentation	20
8.3 Warp test routines	28
8.3.1 Detailed Description	28
8.3.2 Function Documentation	28

9 File Documentation	32
9.1 warp.h File Reference	32
9.1.1 Detailed Description	33
9.1.2 Typedef Documentation	33
9.2 warp_test.h File Reference	33
9.2.1 Detailed Description	34
Index	35

1 Abstract

This package implements an optimal-scaling multigrid solver for the linear systems that arise from the discretization of problems with evolutionary behavior. Typically, solution algorithms for evolution equations are based on a time-marching approach, solving sequentially for one time step after the other. Parallelism in these traditional time-integration techniques is limited to spatial parallelism. However, current trends in computer architectures are leading towards systems with more, but not faster, processors. Therefore, faster compute speeds must come from greater parallelism. One approach to achieve parallelism in time is with multigrid, but extending classical multigrid methods for elliptic operators to this setting is a significant achievement. In this software, we implement a non-intrusive, optimal-scaling time-parallel method based on multigrid reduction techniques. The examples in the package demonstrate optimality of our multigrid-reduction-in-time algorithm (MGRIT) for solving a variety of equations in two and three spatial dimensions. These examples can also be used to show that MGRIT can achieve significant speedup in comparison to sequential time marching on modern architectures.

It is **strongly recommended** that you also read [Parallel Time Integration with Multigrid](#) after reading the [Introduction](#). It is a more in depth discussion of the algorithm and associated experiments.

2 Introduction

2.1 Introduction

2.2 Overview of Warp Algorithm

The goal of Warp is to solve a problem faster than a traditional time marching algorithm. Instead of sequential time marching, Warp solves the problem iteratively by simultaneously updating a space-time solution guess over all time values. The initial solution guess can be anything, even a random function over space-time. The iterative updates to the solution guess are done by constructing a hierarchy of temporal grids, where the finest grid contains all of the time values for the simulation. Each subsequent grid is a coarser grid with fewer time values. The coarsest grid has a trivial number of time steps and can be quickly solved exactly. The effect is that solutions to the time marching problem on the coarser (i.e., cheaper) grids can be used to correct the original finest grid solution. Thus, a problem with many time steps (thousands, tens of thousands or more) can be solved with 10 or 15 Warp iterations, and the overall time to solution can be greatly sped up. However, this is achieved at the cost of more computational resources.

To understand how Warp differs from traditional time marching, consider the simple linear advection equation, $u_t = -cu_x$. The next figure depicts how one would typically evolve a solution here with sequential time stepping. The initial condition is a wave, and this wave propagates sequentially across space as time increases.

Warp instead begins with a solution guess over all of space-time, which for demonstration, we let be random. A Warp iteration then does

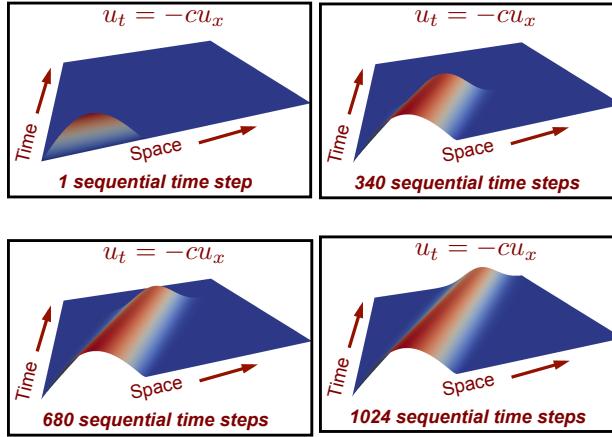


Figure 1: Sequential time stepping.

1. Relaxation on the fine grid, i.e., the grid that contains all of the desired time values
 - Relaxation is just a local application of the time stepping scheme, e.g., backward Euler
2. Restriction to the first coarse grid, i.e., interpolate the problem to a grid that contains fewer time values, say every second or every third time value
3. Relaxation on the first coarse grid
4. Restriction to the second coarse grid and so on...
5. When a coarse grid of trivial size (say 2 time steps) is reached, it is solved exactly.
6. The solution is then interpolated from the coarsest grid to the finest grid

One Warp iteration is called a *cycle* and these cycles continue until the the solution is accurate enough. This is depicted in the next figure, where only a few iterations are required for this simple problem.

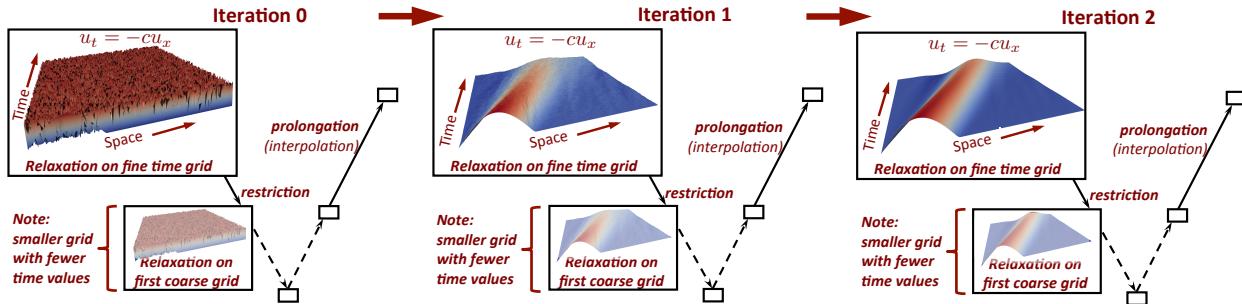


Figure 2: Warp iterations.

There are a few important points to make.

- The coarse time grids allow for global propagation of information across space-time with only one Warp iteration. This is visible in the above figure by observing how the solution is updated from iteration 0 to iteration 1.
- Using coarser (cheaper) grids to correct the fine grid is analogous to spatial multigrid.
- Only a few Warp iterations are required to find the solution over 1024 time steps. Therefore if enough processors are available to parallelize Warp, we can see a speedup over traditional time stepping (more on this later).

- This is a simple example, with evenly space time steps. Warp is structured to handle variable time step sizes and adaptive time step sizes, and these features will be coming.

To firm up our understanding, let's do a little math. Assume that you have a general ODE,

$$u'(t) = f(t, u(t)), \quad u(0) = u_0, \quad t \in [0, T],$$

which you descretize with the one-step integration

$$u_i = \Phi_i(u_{i-1}) + g_i, \quad i = 1, 2, \dots, N.$$

Traditional time marchine would first solve for $i = 1$, then solve for $i = 2$, and so on. This process is equivalent to a forward solve of this system,

$$A\mathbf{u} \equiv \begin{pmatrix} I & & & \\ -\Phi_1 & I & & \\ & \ddots & \ddots & \\ & & -\Phi_N & I \end{pmatrix} \begin{pmatrix} u_0 \\ u_1 \\ \vdots \\ u_N \end{pmatrix} = \begin{pmatrix} g_0 \\ g_1 \\ \vdots \\ g_N \end{pmatrix} \equiv \mathbf{g}$$

or

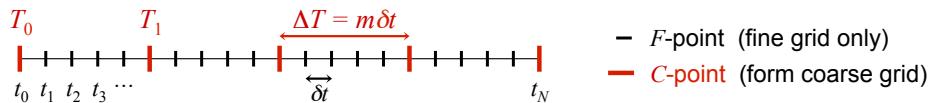
$$A\mathbf{u} = \mathbf{g}.$$

This process is optimal and $O(N)$, but it is sequential. Warp instead solves the system iteratively, with a multigrid reduction method ¹ applied in only the time dimension. This approach is

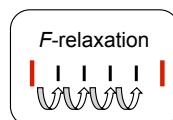
- nonintrusive, in that it coarsens only in time and the user defines Φ
 - Thus, users can continue using existing time stepping codes by wrapping them into our framework.
- optimal and $O(N)$, but $O(N)$ with a higher constant than time stepping
 - Thus with enough computational resources, Warp will outperform sequential time stepping.
- highly parallel

Warp solves this system iteratively by constructing a hierarchy of time grids. We describe the two-grid process, with the multigrid process being a recursive application of the process. We also assume that Φ is constant for notational simplicity.

Warp functions as follows. The next figure depicts a sample timeline of time values, where the time values have been split into C- and F-points. C-points exist on both the fine and coarse time grid, but F-points exist only on the fine time scale. The first task is relaxation and an effective relaxation alternates between C and F sweeps (this is like line-



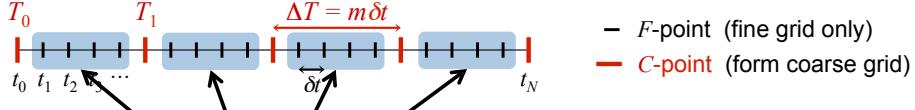
relaxation in space in that the residual is set to 0 for an entire time step). An F sweep simply updates time values by integrating with Φ over all the F-points from one C-point to the next, as depicted next.



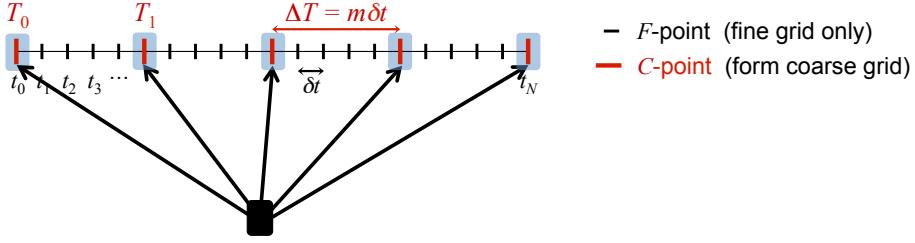
But, such an update can be done simultaneously over all F intervals in parallel, as depicted next.

Following an F sweep we can also do C sweep, as depicted next.

¹ Ries, Manfred, Ulrich Trottenberg, and Gerd Winter. "A note on MGR methods." Linear Algebra and its Applications 49 (1983): 1-26.



- Update all F-points using time propagator Φ



- Update all C-points using time propagator Φ

In general, FCF- and F-relaxation will refer to the relaxation methods used in Warp. We can say

- FCF or F-relaxation is highly parallel.
- But, a sequential component exists equaling the number of F-points between two C-points.
- Warp uses regular coarsening factors, i.e., the spacing of C-points happens every k points.

After relaxation, comes coarse grid correction. The restriction operator R maps fine grid quantities to the coarse grid by simply injecting values at C-points from the fine grid to the coarse grid,

$$R = \begin{pmatrix} I & & & \\ 0 & & & \\ \vdots & & & \\ 0 & & & \\ & I & & \\ & 0 & & \\ & \vdots & & \\ & 0 & & \\ & & & \ddots \end{pmatrix},$$

where the spacing between each I is $m - 1$ block rows. Warp implements an FAS (Full Approximation Scheme) multigrid cycle, and hence the solution guess and residual (i.e., $A, \mathbf{u}, \mathbf{g} - A\mathbf{u}$) are restricted. This is in contrast to linear multigrid which typically restricts only the residual equation to the coarse grid. We choose FAS because it is *nonlinear* multigrid and allows us to solve nonlinear problems. FAS was invented by Achi Brandt, but this [PDF](#) by Van Henson is a good intro.

The main question here is how to form the coarse grid matrix, which in turn asks how to define the coarse grid time stepper Φ_Δ . It is typical to let Φ_Δ simply be Φ but with the coarse time step size $\Delta T = m\delta t$. Thus if

$$A = \begin{pmatrix} I & & & \\ -\Phi & I & & \\ \ddots & \ddots & \ddots & \\ & -\Phi & I \end{pmatrix}$$

then

$$A_\Delta = \begin{pmatrix} I & & & \\ -\Phi_\Delta & I & & \\ & \ddots & \ddots & \\ & & -\Phi_\Delta & I \end{pmatrix},$$

where A_Δ has fewer rows and columns than A , e.g., if we are coarsening in time by 2, A_Δ has one half as many rows and columns. This coarse grid equation

$$A_\Delta \mathbf{v}_\Delta = \mathbf{g}_\Delta$$

is then solved, where the right-hand-side is defined by FAS (see [Two-Grid Algorithm](#)). Finally, FAS defines a coarse grid error approximation \mathbf{e}_Δ , which is interpolated with P_Φ back to the fine grid and added to the current solution guess. Interpolation is equivalent to injecting the coarse grid to the C-points on the fine grid, followed by an F-relaxation sweep. That is,

$$P_\Phi = \begin{pmatrix} I & & & \\ \Phi & & & \\ \Phi^2 & & & \\ \vdots & & & \\ \Phi^{m-1} & & & \\ & I & & \\ & \Phi & & \\ & \Phi^2 & & \\ & \vdots & & \\ & \Phi^{m-1} & & \\ & & \ddots & \end{pmatrix},$$

where m is the coarsening factor.

2.2.1 Two-Grid Algorithm

This two-grid process is captured with this algorithm. Using a recursive coarse grid solve (i.e., step 3 becomes a recursive call) makes the process multilevel. Halting is done based on a residual tolerance. If the operator is linear, this FAS cycle is equivalent to standard linear multigrid. Note that we represent A as a function below, whereas the above notation was simplified for the linear case.

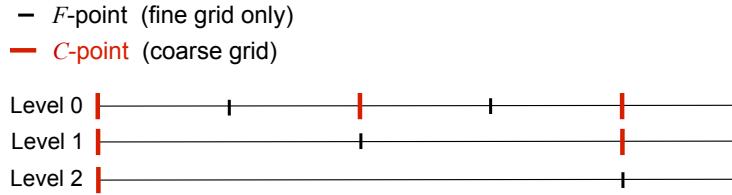
1. Relax on $A(\mathbf{u}) = \mathbf{g}$ using FCF-relaxation
2. Restrict the fine grid approximation and its residual: $\mathbf{u}_\Delta \leftarrow R\mathbf{u}$, $\mathbf{r}_\Delta \leftarrow R(\mathbf{g} - A(\mathbf{u}))$
3. Solve $A_\Delta(\mathbf{v}_\Delta) = A_\Delta(\mathbf{u}_\Delta) + \mathbf{r}_\Delta$
4. Compute the coarse grid error approximation: $\mathbf{e}_\Delta = \mathbf{v}_\Delta - \mathbf{u}_\Delta$
5. Correct: $\mathbf{u} \leftarrow \mathbf{u} + P\mathbf{e}_\Delta$

Caveat: The Warp implementation of FAS differs slightly from standard FAS. In standard FAS, the error is interpolated, using some standard interpolation formula like linear interpolation, to the fine points on the fine grid (here F-points). Instead, given our interpolation operator P_Φ , we add the error to the coarse points on the fine grid (here C-points), and then propagate the *solution* to F-points, like in a reduction method. Thus, F-points are updated in a slightly different, but more exact manner. This strategy allows Warp to save on storage and to not store C-points, while still effectively solving nonlinear problems.

2.2.2 Summary

In summary, a few points are

- Warp is an iterative solver for the global space-time problem.
- The user defines the time stepping routine Φ and can wrap existing code to accomplish this.
- Warp convergence will depend heavily on how well Φ_Δ approximates Φ^m , that is how well a time step size of $m\delta t = \Delta T$ will approximate m applications of the same time integrator for a time step size of δt . This is a subject of research, but this approximation need not capture fine scale behavior, which is instead captured by relaxation on the fine grid.
- The coarsest grid is solved exactly, i.e., sequentially, which can be a bottleneck for two-level methods like Parareal,² but not for a multilevel scheme like Warp where the coarsest grid is of trivial size.
- By forming the coarse grid to have the same sparsity structure and time stepper as the fine grid, the algorithm can recur easily and efficiently.
- Interpolation is ideal or exact, in that an application of interpolation leaves a zero residual at all F-points.
- The process is applied recursively until a trivially sized temporal grid is reached, e.g., 2 or 3 time points. Thus, the coarsening rate m determines how many levels there are in the hierarchy. For instance in this figure, a 3 level hierarchy is shown. Three levels are chosen because there are six time points, $m = 2$ and $m^2 < 6 \leq m^3$. If the coarsening rate had been $m = 4$ then there would only be two levels because, there would be no more points to coarsen!



By default, Warp will subdivide the time domain into evenly sized time steps. Warp is structured to handle variable time step sizes and adaptive time step sizes, and these features are coming.

2.3 Overview of Warp Code

Warp is designed to run in conjunction with an existing application code that can be wrapped per our interface. This application code will implement some time marching type simulation like fluid flow. Essentially, the user has to take their application code and extract a stand-alone time-stepping function Φ that can evolve a solution from one time value to another, regardless of time step size. After this is done, the Warp code takes care of the parallelism in the time dimension.

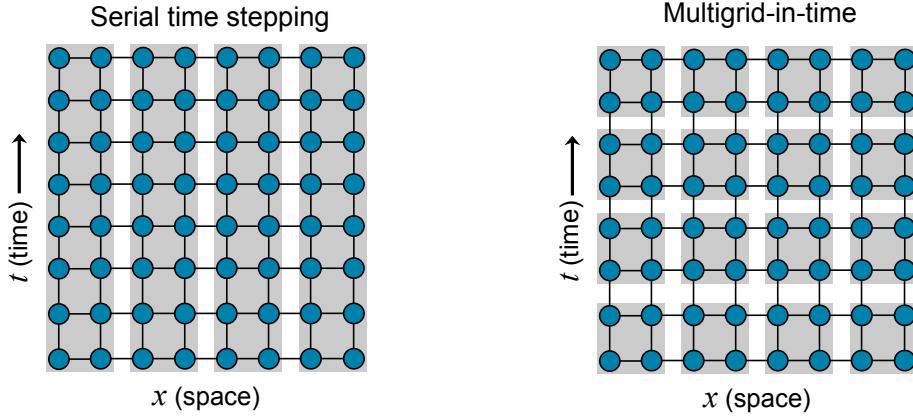
Warp

- is written in C and can easily interface with Fortran and C++
- uses MPI for parallelism
- self documents through comments in the source code and through *.md files

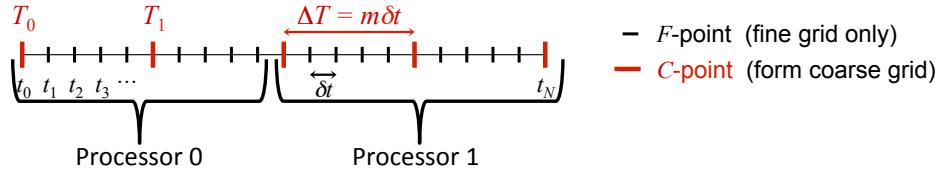
2.3.1 Parallel decomposition and memory

- Warp decomposes the problem in parallel as depicted next. As you can see, traditional time stepping only stores one time step at a time, but only enjoys a spatial data decomposition and spatial parallelism. On the other hand, Warp stores multiple time steps simultaneously and each processor holds a space-time chunk reflecting both the spatial and temporal parallelism.

² Lions, J., Yvon Maday, and Gabriel Turinici. "A"parareal"in time discretization of PDE's." Comptes Rendus de l'Academie des Sciences Series I Mathematics 332.7 (2001): 661-668.



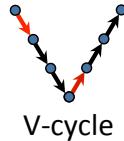
- Warp only handles temporal parallelism and is agnostic to the spatial decomposition. See [warp_SplitCommworld](#). Each processor owns a certain number of CF intervals of points, as depicted next, where each processor owns 2 CF intervals. Warp distributes Intervals evenly on the finest grid.



- Storage is greatly minimized by only storing C-points. Whenever an F-point is needed, it is generated by F-relaxation. That is, we only store the red C-point time values in the previous figure. Coarsening can be aggressive with $m = 8, 16, 32$, so the storage requirements of Warp are significantly reduced when compared to storing all of the time values.
By only storing data at C-points, we effect a subtle change to the standard FAS algorithm (see [Two-Grid Algorithm](#)).
- In practice, storing only one space-time slab is advisable. That is, solve for as many time steps (say k time steps) as you have available memory for. Then move on to the next k time steps.

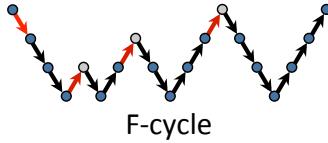
2.3.2 Cycling and relaxation strategies

There are two main cycling strategies available in Warp, F-and V-cycles. These two cycles differ in how often and the order in which coarse levels are visited. A V-cycle is depicted next, and is a simple recursive application of the [Two-Grid Algorithm](#).



An F-cycle visits coarse grids more frequently and in a different order. Essentially, an F-cycle uses a V-cycle as the post-smoother, which is an expensive choice for relaxation. But, this extra work gives you a closer approximation to a two-grid cycle, and a faster convergence rate at the extra expense of more work. The effectiveness of a V-cycle as a relaxation scheme can be seen in Figure 2, where one V-cycle globally propagates and *smoothes* the error. The cycling strategy of an F-cycle is depicted next.

Next, we make a few points about F- versus V-cycles.



- One V-cycle iteration is cheaper than one F-cycle iteration.
- But, F-cycles often converge more quickly. For some test cases, this difference can be quite large. The cycle choice for the best time to solution will be problem dependent. See [Heat equation example](#) for a case study of cycling strategies.

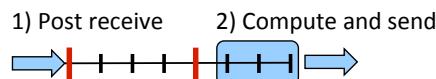
The number of FC relaxation sweeps is another important algorithmic setting. Note that at least one F-relaxation sweep is always done on a level. A few summary points about relaxation are as follows.

- Using FCF (or even FCFCF, FCFCFCF) relaxation, corresponding to passing `warp_SetNRelax` a value of 1, 2 or 3 respectively, will result in a Warp cycle that converges more quickly as the number of relaxations grows.
- But as the number of relaxations grows, each Warp cycle becomes more expensive. The optimal relaxation strategy for the best time to solution will be problem dependent.
- However, a good first step is to try FCF on all levels (i.e., `warp_SetNRelax(core, -1, 1)`).
- A common optimization is to first set FCF on all levels (i.e., `warp_setnrelax(core, -1, 1)`), but then overwrite the FCF option on level 0 so that only F-relaxation is done on level 0, (i.e., `warp_setnrelax(core, 0, 1)`). This strategy can work well with F-cycles.
- See [Heat equation example](#) for a case study of relaxation strategies.

Last, [Parallel Time Integration with Multigrid](#) has a more in depth case study of cycling and relaxation strategies

2.3.3 Overlapping communication and computation

Warp effectively overlaps communication and computation. The main computational kernel of Warp is relaxation (C or F). At the start of each sweep, each processor first posts a send at its left-most point, and then carries out F-relaxation on its right-most interval in order to send the next processor the data that it needs. If each processor has multiple intervals at this Warp level, this should allow for complete overlap.



2.3.4 Configuring the Warp Hierarchy

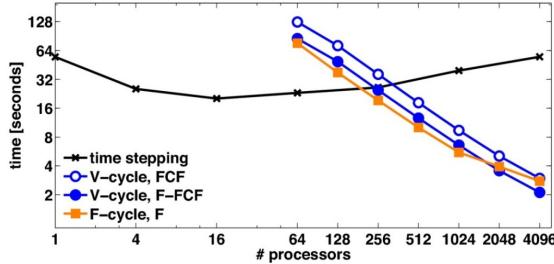
Some of the more basic Warp function calls allow you to control aspects discussed here.

- `warp_SetFMG`: switches between using F- and V-cycles.
- `warp_SetMaxIter`: sets the maximum number of Warp iterations
- `warp_SetCFactor`: sets the coarsening factor for any (or all levels)
- `warp_SetNRelax`: sets the number of CF-relaxation sweeps for any (or all levels)
- `warp_SetRelTol`, `warp_SetAbsTol`: sets the stopping tolerance

- `warp_SetMaxCoarse`: sets the maximum coarse grid size, in terms of C-points
- `warp_SetMaxLevels`: sets the maximum number of levels in the Warp hierarchy

2.3.5 Heat equation example

Here is some experimental data for the 2D heat equation, $u_t = u_{xx} + u_{yy}$ generated by examples/drive-02. The problem



setup is as follows.

- Backwards Euler is used as the time stepper.
- We used a Linux cluster with 4 cores per node, a Sandybridge Intel chipset, and a fast Infiniband interconnect.
- The space-time problem size was $129^2 \times 16,192$ over the unit cube $[0, 1] \times [0, 1] \times [0, 1]$.
- The coarsening factor was $m = 16$ on the finest level and $m = 2$ on coarser levels.
- Since 16 processors optimized the serial time stepping approach, 16 processors in space are also used for the Warp experiments. So for instance 512 processors in the plot corresponds to 16 processors in space and 32 processors in time, $16 * 32 = 512$. Thus, each processor owns a space-time hypercube of $(129^2/16) \times (16,192/32)$. See [Parallel decomposition and memory](#) for a depiction of how Warp breaks the problem up.
- Various relaxation and V and F cycling strategies are experimented with.
 - *V-cycle, FCF* denotes V-cycles and FCF-relaxation on each level.
 - *V-cycle, F-FCF* denotes V-cycles and F-relaxation on the finest level and FCF-relaxation on all coarser levels.
 - *F-cycle, F* denotes F-cycles and F-relaxation on each level.
- The initial guess at time values for $t > 0$ is zero, which is typical.

Regarding the performance, we can say

- The best speedup is 10x and this would grow if more processors were available.
- Although not shown, the iteration counts here are about 10-15 Warp iterations. See [Parallel Time Integration with Multigrid](#) for the exact iteration counts.
- At smaller core counts, serial time stepping is faster. But at about 256 processors, there is a crossover and Warp is faster.
- You can see the impact of the cycling and relaxation strategies discussed in [Cycling and relaxation strategies](#). For instance, even though *V-cycle, F-FCF* is a weaker relaxation strategy than *V-cycle, FCF* (i.e., the Warp convergence is slower), *V-cycle, F-FCF* has a faster time to solution than *V-cycle, FCF* because each cycle is cheaper.
- In general, one level of aggressive coarsening (here by a factor 16) followed by slower coarsening was found to be best on this machine.

Achieving the best speedup can require some tuning, and it is recommended to read [Parallel Time Integration with Multigrid](#) where this 2D heat equation example is explored in much more detail.

2.4 Summary of Warp

- Warp applies multigrid to the time dimension.
 - This exposes concurrency in the time dimension.
 - The potential for speedup is large, 10x, 100x, ...
- This is a non-intrusive approach, with an unchanged time discretization defined by user.
- Parallel time integration is only useful beyond some scale. This is evidenced by the experimental results below. For smaller numbers of cores sequential time stepping is faster, but at larger core counts Warp is much faster.
- The more time steps that you can parallelize over, the better your speedup will be.
- Warp is optimal for a variety of parabolic problems (see the examples directory).

3 Example

A Simple Example

User Defined Structures and Wrappers

As mentioned, the user must wrap their existing time stepping routine per the Warp interface. To do this, the user must define two data structures and some wrapper routines. To make the idea more concrete, we now give these function definitions from examples/drive-01, which implements a scalar ODE, $u_t = \lambda u$.

The two data structures are:

1. **App:** This holds a wide variety of information and is *global* in that it is passed to every function. This structure holds everything that the user will need to carry out a simulation. Here, this is just the global MPI communicator and few values describing the temporal domain.

```
typedef struct _warp_App_struct
{
    MPI_Comm    comm;
    double      tstart;
    double      tstop;
    int         ntime;

} my_App;
```

2. **Vector:** this defines (roughly) a state vector at a certain time value. It could also contain any other information related to this vector which is needed to evolve the vector to the next time value, like mesh information. Here, the vector is just a scalar double.

```
typedef struct _warp_Vector_struct
{
    double      value;

} my_Vector;
```

The user must also define a few wrapper routines. Note, that the app structure is the first argument to every function.

1. **Phi:** This function tells Warp how to take a time step, and is the core user routine. The user must advance the vector u from time $tstart$ to time $tstop$. Here advancing the solution just involves the scalar λ . The *rfactor_ptr* and *accuracy* parameters are advanced topics not used here.

Importantly, the g_i function (from [Overview of Warp Algorithm](#)) must be incorporated into Phi, so that $\Phi(u_i) \rightarrow u_{i+1}$

```

int
my_Phi(warp_App      app,
       double       tstart,
       double       tstop,
       double       accuracy,
       warp_Vector u,
       int          *rfactor_ptr)
{
    /* On the finest grid, each value is half the previous value */
    (u->value) = pow(0.5, tstop-tstart)*(u->value);

    /* Zero rhs for now */
    (u->value) += 0.0;

    /* no refinement */
    *rfactor_ptr = 1;

    return 0;
}

```

2. **Init:** This function tells Warp how to initialize a vector at time t . Here that is just allocating and setting a scalar on the heap.

```

int
my_Init(warp_App      app,
        double       t,
        warp_Vector *u_ptr)
{
    my_Vector *u;

    u = (my_Vector *) malloc(sizeof(my_Vector));
    if (t == 0.0)
    {
        /* Initial guess */
        (u->value) = 1.0;
    }
    else
    {
        /* Random between 0 and 1 */
        (u->value) = ((double)rand()) / RAND_MAX;
    }
    *u_ptr = u;

    return 0;
}

```

3. **Clone:** This function tells Warp how to clone a vector into a new vector.

```

int
my_Clone(warp_App      app,
          warp_Vector u,
          warp_Vector *v_ptr)
{
    my_Vector *v;

    v = (my_Vector *) malloc(sizeof(my_Vector));
    (v->value) = (u->value);
    *v_ptr = v;

    return 0;
}

```

4. **Free:** This function tells Warp how to free a vector.

```

int
my_Free(warp_App      app,
        warp_Vector u)
{
    free(u);

    return 0;
}

```

5. **Sum:** This function tells Warp how to sum two vectors (AXPY operation).

```
int
my_Sum(warp_App    app,
       double      alpha,
       warp_Vector x,
       double      beta,
       warp_Vector y)
{
    (y->value) = alpha*(x->value) + beta*(y->value);

    return 0;
}
```

6. **Dot:** This function tells Warp how to take the dot product of two vectors.

```
int
my_Dot(warp_App    app,
        warp_Vector u,
        warp_Vector v,
        double      *dot_ptr)
{
    double dot;

    dot = (u->value)*(v->value);
    *dot_ptr = dot;

    return 0;
}
```

7. **Write:** This function tells Warp how to write a vector at time t to screen, file, etc... The user defines what is appropriate output. Notice how you are told the time value of the vector u and even more information in $status$. This lets you tailor the output to only certain time values.

If `write_level` is 2 (see [warp_SetWriteLevel](#)), then `Write` is called every Warp iteration and on every Warp level. In this case, `status` can be queried using the `warp_Get**Status()` functions, to determine the current Warp level and iteration. This allows for even more detailed tracking of the simulation.

See examples/drive-02 and examples/drive-04 for more advanced uses of the Write function. Drive-04 writes to a GLVIS visualization port, and examples/drive-02 writes to .vtu files.

```
int
my_Write(warp_App    app,
          double      t,
          warp_Status status,
          warp_Vector u)
{
    MPI_Comm    comm    = (app->comm);
    double     tstart = (app->tstart);
    double     tstop  = (app->tstop);
    int       ntime   = (app->ntime);
    int       index, myid;
    char      filename[255];
    FILE     *file;

    index = ((t-tstart) / ((tstop-tstart)/ntime) + 0.1);

    MPI_Comm_rank(comm, &myid);

    sprintf(filename, "%s.%07d.%05d", "drive-01.out", index, myid);
    file = fopen(filename, "w");
    fprintf(file, "%1.14e\n", (u->value));
    fflush(file);
    fclose(file);

    return 0;
}
```

8. **BufSize, BufPack, BufUnpack:** These three routines tell Warp how to communicate vectors between processors. `BufPack` packs a vector into a `void *` buffer for MPI and then `BufUnPack` unpacks it from `void *` to vector. Here doing that for a scalar is trivial. `BufSize` computes the upper bound for the size of an arbitrary vector.

```
int
my_BufSize(warp_App    app,
```

```

        int      *size_ptr)
{
    *size_ptr = sizeof(double);
    return 0;
}

int
my_BufPack(warp_App      app,
            warp_Vector  u,
            void         *buffer)
{
    double *dbufbuffer = buffer;

    dbuffer[0] = (u->value);

    return 0;
}

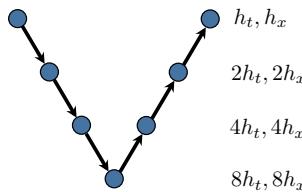
int
my_BufUnpack(warp_App      app,
              void         *buffer,
              warp_Vector *u_ptr)
{
    double  *dbufbuffer = buffer;
    my_Vector *u;

    u = (my_Vector *) malloc(sizeof(my_Vector));
    (u->value) = dbuffer[0];
    *u_ptr = u;

    return 0;
}

```

9. **Coarsen, Restrict** (optional): These are advanced options that allow for coarsening in space while you coarsen in time. This is useful for maintaining stable explicit schemes on coarse time scales and is not needed here. See for instance examples/drive-04 and examples/drive-05 which use these routines.
These functions allow you vary the spatial mesh size on Warp levels as depicted here where the spatial and temporal grid sizes are halved every level.



10. Adaptive and variable time stepping is in the works to be implemented. The *rfactor* parameter in *Phi* will allow this.

Running Warp

A typical flow of events in the `main` function is to first initialize the `app` structure.

```

/* set up app structure */
app = (my_App *) malloc(sizeof(my_App));
(app->comm)   = comm;
(app->tstart) = tstart;
(app->tstop)  = tstop;
(app->ntime)  = ntime;

```

Then, the data structure definitions and wrapper routines are passed to Warp. The core structure is used by Warp for internal data structures.

```
warp_Core core;
warp_Init(MPI_COMM_WORLD, comm, tstart, tstop, ntime, app,
           my_Phi, my_Init, my_Clone, my_Free, my_Sum, my_Dot, my_Write,
           my_BufSize, my_BufPack, my_BufUnpack,
           &core);
```

Then, Warp options are set.

```
warp_SetPrintLevel( core, 1);
warp_SetMaxLevels(core, max_levels);
warp_SetNRelax(core, -1, nrelax);
warp_SetAbsTol(core, tol);
warp_SetCFactor(core, -1, cfactor);
warp_SetMaxIter(core, max_iter);
```

Then, the simulation is run.

```
warp_Drive(core);
```

Then, we clean up.

```
warp_Destroy(core);
```

Finally, to run drive-01, type

```
drive-01 -ml 5
```

This will run drive-01. See examples/drive-0* for more extensive examples.

Testing Warp

The best overall test for Warp, is to set the maximum number of levels to 1 (see [warp_SetMaxLevels](#)) which will carry out a sequential time stepping test. Take the output given to you by your *Write* function and compare it to output from a non-Warp run. Is everything OK? Once this is complete, repeat for multilevel Warp, and check that the solution is correct (that is, it matches a serial run to within tolerance).

At a lower level, to do sanity checks of your data structures and wrapper routines, there are also Warp test functions, which can be easily run. The test routines also take as arguments the app structure, spatial communicator *comm_x*, a stream like *stdout* for test output and a time step size to test *dt*. After these arguments, function pointers to wrapper routines are the rest of the arguments. Some of the tests can return a boolean variable to indicate correctness.

```
/* Test init(), write(), free() */
warp_TestInitWrite( app, comm_x, stdout, dt, my_Init, my_Write, my_Free);

/* Test clone() */
warp_TestClone( app, comm_x, stdout, dt, my_Init, my_Write, my_Free, my_Clone);

/* Test sum() */
warp_TestSum( app, comm_x, stdout, dt, my_Init, my_Write, my_Free, my_Clone, my_Sum);

/* Test dot() */
correct = warp_TestDot( app, comm_x, stdout, dt, my_Init, my_Free, my_Clone, my_Sum, my_Dot);

/* Test bufsize(), bufpack(), bufunpack() */
correct = warp_TestBuf( app, comm_x, stdout, dt, my_Init, my_Free, my_Sum, my_Dot, my_BufSize, my_BufPack, my_BufUnpack);

/* Test coarsen and refine */
correct = warp_TestCoarsenRefine(app, comm_x, stdout, 0.0, dt, 2*dt, my_Init,
                                 my_Write, my_Free, my_Clone, my_Sum, my_Dot, my_CoarsenInjection,
                                 my_Refine);
correct = warp_TestCoarsenRefine(app, comm_x, stdout, 0.0, dt, 2*dt, my_Init,
                                 my_Write, my_Free, my_Clone, my_Sum, my_Dot, my_CoarsenBilinear,
                                 my_Refine);
```

4 Building Warp

- To specify the compilers, flags and options for your machine, edit makefile.inc. For now, we keep it simple and avoid using configure or cmake.
- To make the library, libwarp.a,
`$ make`
- To make the examples
`$ make all`
- The makefile lets you pass some parameters like debug with
`$ make debug=yes`
or
`$ make all debug=yes`
It would also be easy to add additional parameters, e.g., to compile with insure.
- To set compilers and library locations, look in makefile.inc where you can set up an option for your machine to define simple stuff like
`CC = mpicc`
`MPICC = mpicc`
`MPICXX = mpiCC`
`LFLAGS = -lm`

5 Compiling and running the examples

Type

```
drive-0* -help
```

for instructions on how to run any driver.

To run the examples, type

```
mpirun -np 4 drive-* [args]
```

1. drive-01 is the simplest example. It implements a scalar ODE and can be compiled and run with no outside dependencies.
2. drive-02 implements the 2D heat equation on a regular grid. You must have `hypre` installed and these variables in examples/Makefile set correctly
`HYPRE_DIR = ../../linear_solvers/hypre`
`HYPRE_FLAGS = -I$(HYPRE_DIR)/include`
`HYPRE_LIB = -L$(HYPRE_DIR)/lib -lHYPRE`
3. drive-03 implements the 3D heat equation on a regular grid, and assumes `hypre` is installed just like drive-02.
4. drive-05 implements the 2D heat equation on a regular grid, but it uses spatial coarsening. This allows you to use explicit time stepping on each Warp level, regardless of time step size. It assumes `hypre` is installed just like drive-02.
5. drive-04 is a sophisticated test bed for various PDEs, mostly parabolic. It relies on the `mfem` package to create general finite element discretizations for the spatial problem. Other packages must be installed in this order.

- Unpack and install `Metis`
- Unpack and install `hypre`
- Unpack and install `mfem`. Make the serial version of mfem first by only typing `make`. Then make sure to set these variables correctly in the mfem Makefile:
`USE_METIS_5 = YES`
`HYPRE_DIR = where_ever_linear_solvers_is/hypre`
- Make `GLVIS`, which needs serial mfem. Set these variables in the glvis makefile
`MFEM_DIR = mfem_location`
`MFEM_LIB = -L$(MFEM_DIR) -lmfem`

- Go back to the mfem directory and type
`make clean`
`make parallel`
- Go to warp/examples and set these Makefile variables,
`METIS_DIR = ../../metis-5.1.0/lib`
`MFEM_DIR = ../../mfem`
`MFEM_FLAGS = -I$(MFEM_DIR)`
`MFEM_LIB = -L$(MFEM_DIR) -lmfem -L$(METIS_DIR) -lmetis`
 then type
`make drive-04`
- To run `drive-04` and `glvis`, open two windows. In one, start a `glvis` session
`./glvis`
 Then, in the other window, run `drive-04`
`mpirun -np ... drive-04 [args]`
`Glvis will listen on a port to which drive-04 will dump visualization information.`

6 Module Index

6.1 Modules

Here is a list of all modules:

User-written routines	16
User interface routines	19
Warp test routines	28

7 File Index

7.1 File List

Here is a list of all files with brief descriptions:

warp.h Define headers for user interface routines	32
warp_test.h Define headers for Warp test routines	33

8 Module Documentation

8.1 User-written routines

TypeDefs

- `typedef struct _warp_App_struct * warp_App`
- `typedef struct _warp_Vector_struct * warp_Vector`
- `typedef struct _warp_Status_struct * warp_Status`

- `typedef warp_Int(* warp_PtFcnPhi)(warp_App app, warp_Real tstart, warp_Real tstop, warp_Real accuracy, warp_Vector u, warp_Int *rfactor_ptr)`
- `typedef warp_Int(* warp_PtFcnInit)(warp_App app, warp_Real t, warp_Vector *u_ptr)`
- `typedef warp_Int(* warp_PtFcnClone)(warp_App app, warp_Vector u, warp_Vector *v_ptr)`
- `typedef warp_Int(* warp_PtFcnFree)(warp_App app, warp_Vector u)`
- `typedef warp_Int(* warp_PtFcnSum)(warp_App app, warp_Real alpha, warp_Vector x, warp_Real beta, warp_Vector y)`
- `typedef warp_Int(* warp_PtFcnDot)(warp_App app, warp_Vector u, warp_Vector v, warp_Real *dot_ptr)`
- `typedef warp_Int(* warp_PtFcnWrite)(warp_App app, warp_Real t, warp_Status status, warp_Vector u)`
- `typedef warp_Int(* warp_PtFcnBufSize)(warp_App app, warp_Int *size_ptr)`
- `typedef warp_Int(* warp_PtFcnBufPack)(warp_App app, warp_Vector u, void *buffer)`
- `typedef warp_Int(* warp_PtFcnBufUnpack)(warp_App app, void *buffer, warp_Vector *u_ptr)`
- `typedef warp_Int(* warp_PtFcnCoarsen)(warp_App app, warp_Real tstart, warp_Real f_tminus, warp_Real f_tplus, warp_Real c_tminus, warp_Real c_tplus, warp_Vector fu, warp_Vector *cu_ptr)`
- `typedef warp_Int(* warp_PtFcnRefine)(warp_App app, warp_Real tstart, warp_Real f_tminus, warp_Real f_tplus, warp_Real c_tminus, warp_Real c_tplus, warp_Vector cu, warp_Vector *fu_ptr)`

8.1.1 Detailed Description

These are all user-written data structures and routines. There are two data structures (`warp_App` and `warp_Vector`) for the user to define. And, there are a variety of function interfaces (defined through function pointer declarations) that the user must implement.

8.1.2 Typedef Documentation

8.1.2.1 `typedef struct _warp_App_struct* warp_App`

This holds a wide variety of information and is `global` in that it is passed to every function. This structure holds everything that the user will need to carry out a simulation. For a simple example, this could just hold the global MPI communicator and a few values describing the temporal domain.

8.1.2.2 `typedef warp_Int(* warp_PtFcnBufPack)(warp_App app,warp_Vector u,void *buffer)`

This allows warp to send messages containing warp_Vectors. This routine packs a vector u into a `void * buffer` for MPI.

8.1.2.3 `typedef warp_Int(* warp_PtFcnBufSize)(warp_App app,warp_Int *size_ptr)`

This routine tells Warp message sizes by computing an upper bound in bytes for an arbitrary `warp_Vector`. This size must be an upper bound for what BufPack and BufUnPack will assume.

8.1.2.4 `typedef warp_Int(* warp_PtFcnBufUnpack)(warp_App app,void *buffer,warp_Vector *u_ptr)`

This allows warp to receive messages containing warp_Vectors. This routine unpacks a `void * buffer` from MPI into a `warp_Vector`.

8.1.2.5 `typedef warp_Int(* warp_PtFcnClone)(warp_App app,warp_Vector u,warp_Vector *v_ptr)`

Clone u into v_ptr

8.1.2.6 `typedef warp_Int(* warp_PtFcnCoarsen)(warp_App app,warp_Real tstart,warp_Real f_tminus,warp_Real f_tplus,warp_Real c_tminus,warp_Real c_tplus,warp_Vector fu,warp_Vector *cu_ptr)`

spatial coarsening (optional). Allows the user to coarsen when going from a fine time grid to a coarse time grid. This function is called on every vector at each level, thus you can coarsen the entire space time domain. The action of this function should match the [warp_PtFcnRefine](#) function.

8.1.2.7 `typedef warp_Int(* warp_PtFcnDot)(warp_App app,warp_Vector u,warp_Vector v,warp_Real *dot_ptr)`

Carry out a dot product $dot_ptr = \langle u, v \rangle$

8.1.2.8 `typedef warp_Int(* warp_PtFcnFree)(warp_App app,warp_Vector u)`

Free and deallocate u

8.1.2.9 `typedef warp_Int(* warp_PtFcnInit)(warp_App app,warp_Real t,warp_Vector *u_ptr)`

Initializes a vector u_ptr at time t

8.1.2.10 `typedef warp_Int(* warp_PtFcnPhi)(warp_App app,warp_Real tstart,warp_Real tstop,warp_Real accuracy,warp_Vector u,warp_Int *rfactor_ptr)`

Defines the central time stepping function that the user must write. The user must advance the vector u from time $tstart$ to time $tstop$. The $rfactor_ptr$ and $accuracy$ inputs are advanced topics. $rfactor_ptr$ allows the user to tell Warp to refine this time interval.

8.1.2.11 `typedef warp_Int(* warp_PtFcnRefine)(warp_App app,warp_Real tstart,warp_Real f_tminus,warp_Real f_tplus,warp_Real c_tminus,warp_Real c_tplus,warp_Vector cu,warp_Vector *fu_ptr)`

spatial refinement (optional). Allows the user to refine when going from a coarse time grid to a fine time grid. This function is called on every vector at each level, thus you can refine the entire space time domain. The action of this function should match the [warp_PtFcnCoarsen](#) function.

8.1.2.12 `typedef warp_Int(* warp_PtFcnSum)(warp_App app,warp_Real alpha,warp_Vector x,warp_Real beta,warp_Vector y)`

AXPY, $\alpha x + \beta y \rightarrow y$

8.1.2.13 `typedef warp_Int(* warp_PtFcnWrite)(warp_App app,warp_Real t,warp_Status status,warp_Vector u)`

Write the vector u at time t to screen, file, etc... The user decides what is appropriate. Notice how you are told the time value of the vector u and even more information in $status$. This lets you tailor the output to only certain time values.

If `write_level` is 2 (see [warp_SetWriteLevel](#)), then `write` is called every Warp iteration and on every Warp level. In this case, `status` can be queried using the `warp_Get**Status()` functions, to determine the current warp level and iteration. This allows for even more detailed tracking of the simulation.

8.1.2.14 `typedef struct _warp_Status_struct* warp_Status`

Points to the status structure defined in `_warp.h` This is NOT a user-defined structure.

8.1.2.15 `typedef struct _warp_Vector_struct* warp_Vector`

This defines (roughly) a state vector at a certain time value. It could also contain any other information related to this vector which is needed to evolve the vector to the next time value, like mesh information.

8.2 User interface routines

TypeDefs

- `typedef struct _warp_Core_struct * warp_Core`

Functions

- `warp_Int warp_Init (MPI_Comm comm_world, MPI_Comm comm, warp_Real tstart, warp_Real tstop, warp_Int ntime, warp_App app, warp_PtFcnPhi phi, warp_PtFcnInit init, warp_PtFcnClone clone, warp_PtFcnFree free, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnWrite write, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack, warp_Core *core_ptr)`
- `warp_Int warp_Drive (warp_Core core)`
- `warp_Int warp_Destroy (warp_Core core)`
- `warp_Int warp_PrintStats (warp_Core core)`
- `warp_Int warp_SetLoosexTol (warp_Core core, warp_Int level, warp_Real loose_tol)`
- `warp_Int warp_SetTightxTol (warp_Core core, warp_Int level, warp_Real tight_tol)`
- `warp_Int warp_SetMaxLevels (warp_Core core, warp_Int max_levels)`
- `warp_Int warp_SetMaxCoarse (warp_Core core, warp_Int max_coarse)`
- `warp_Int warp_SetAbsTol (warp_Core core, warp_Real atol)`
- `warp_Int warp_SetRelTol (warp_Core core, warp_Real rtol)`
- `warp_Int warp_SetNRelax (warp_Core core, warp_Int level, warp_Int nrelax)`
- `warp_Int warp_SetCFactor (warp_Core core, warp_Int level, warp_Int cfactor)`
- `warp_Int warp_SetMaxIter (warp_Core core, warp_Int max_iter)`
- `warp_Int warp_SetFMG (warp_Core core)`
- `warp_Int warp_SetNFMGVcyc (warp_Core core, warp_Int nfmvg_Vcyc)`
- `warp_Int warp_SetSpatialCoarsen (warp_Core core, warp_PtFcnCoarsen coarsen)`
- `warp_Int warp_SetSpatialRefine (warp_Core core, warp_PtFcnRefine refine)`
- `warp_Int warp_SetPrintLevel (warp_Core core, warp_Int print_level)`
- `warp_Int warp_SetPrintFile (warp_Core core, const char *printfile_name)`
- `warp_Int warp_SetWriteLevel (warp_Core core, warp_Int write_level)`
- `warp_Int warp_SplitCommworld (const MPI_Comm *comm_world, warp_Int px, MPI_Comm *comm_x, MPI_Comm *comm_t)`
- `warp_Int warp_GetStatusResidual (warp_Status status, warp_Real *rnorm_ptr)`
- `warp_Int warp_GetStatusIter (warp_Status status, warp_Int *iter_ptr)`
- `warp_Int warp_GetStatusLevel (warp_Status status, warp_Int *level_ptr)`
- `warp_Int warp_GetStatusDone (warp_Status status, warp_Int *done_ptr)`
- `warp_Int warp_GetNumIter (warp_Core core, warp_Int *niter_ptr)`
- `warp_Int warp_GetRNorm (warp_Core core, warp_Real *rnorm_ptr)`

8.2.1 Detailed Description

these are interface routines to initialize and run Warp

8.2.2 Typedef Documentation

8.2.2.1 `typedef struct _warp_Core_struct* warp_Core`

points to the core structure defined in `_warp.h`

8.2.3 Function Documentation

8.2.3.1 warp_Int warp_Destroy (warp_Core *core*)

Clean up and destroy core.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
-------------	-------------------------------

8.2.3.2 warp_Int warp_Drive(warp_Core core)

Carry out a simulation with Warp. Integrate in time.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
-------------	-------------------------------

8.2.3.3 warp_Int warp_GetNumIter(warp_Core core, warp_Int * niter_ptr)

After Drive() finishes, this returns the number of iterations taken.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>niter_ptr</i>	output, holds number of iterations taken

8.2.3.4 warp_Int warp_GetRNorm(warp_Core core, warp_Real * rnorm_ptr)

After Drive() finishes, this returns the last measured residual norm.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>rnorm_ptr</i>	output, holds final residual norm

8.2.3.5 warp_Int warp_GetStatusDone(warp_Status status, warp_Int * done_ptr)

Return whether warp is done for the current status object

done_ptr = 1 indicates that this is the last call to Write, because Warp has stopped iterating (either maxiter has been reached, or the tolerance has been met).

Parameters

<i>status</i>	structure containing current simulation info
<i>done_ptr</i>	output, =1 if warp has finished and this is the final Write, else =0

8.2.3.6 warp_Int warp_GetStatusIter(warp_Status status, warp_Int * iter_ptr)

Return the iteration for the current status object.

Parameters

<i>status</i>	structure containing current simulation info
<i>iter_ptr</i>	output, current iteration number

8.2.3.7 warp_Int warp_GetStatusLevel(warp_Status status, warp_Int * level_ptr)

Return the warp level for the current status object.

Parameters

<i>status</i>	structure containing current simulation info
<i>level_ptr</i>	output, current level in Warp

8.2.3.8 warp_Int warp_GetStatusResidual (warp_Status *status*, warp_Real * *rnorm_ptr*)

Return the residual for the current status object.

Parameters

<i>status</i>	structure containing current simulation info
<i>rnorm_ptr</i>	output, current residual norm

8.2.3.9 warp_Int warp_Init (MPI_Comm *comm_world*, MPI_Comm *comm*, warp_Real *tstart*, warp_Real *tstop*, warp_Int *ntime*, warp_App *app*, warp_PtFcnPhi *phi*, warp_PtFcnInit *init*, warp_PtFcnClone *clone*, warp_PtFcnFree *free*, warp_PtFcnSum *sum*, warp_PtFcnDot *dot*, warp_PtFcnWrite *write*, warp_PtFcnBufSize *bufsize*, warp_PtFcnBufPack *bufpack*, warp_PtFcnBufUnpack *bufunpack*, warp_Core * *core_ptr*)

Create a core object with the required initial data.

This core is used by Warp for internal data structures. The output is *core_ptr* which points to the newly created warp_Core structure.

Parameters

<i>comm_world</i>	Global communicator for space and time
<i>comm</i>	Communicator for temporal dimension
<i>tstart</i>	start time
<i>tstop</i>	End time
<i>ntime</i>	Initial number of temporal grid values
<i>app</i>	User-defined _warp_App structure
<i>phi</i>	User time stepping routine to advance a warp_Vector forward one step
<i>init</i>	Initialize a warp_Vector on the finest temporal grid
<i>clone</i>	Clone a warp_Vector
<i>free</i>	Free a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors
<i>dot</i>	Compute dot product between two warp_Vectors
<i>write</i>	Writes a warp_Vector to file, screen
<i>bufsize</i>	Computes size for MPI buffer for one warp_Vector
<i>bufpack</i>	Packs MPI buffer to contain one warp_Vector
<i>bufunpack</i>	Unpacks MPI buffer into a warp_Vector
<i>core_ptr</i>	Pointer to warp_Core (_warp_Core) struct

8.2.3.10 warp_Int warp_PrintStats (warp_Core *core*)

Print statistics after a Warp run.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
-------------	-------------------------------

8.2.3.11 warp_Int warp_SetAbsTol (warp_Core *core*, warp_Real *atol*)

Set absolute stopping tolerance.

Recommended option over relative tolerance

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>atol</i>	absolute stopping tolerance

8.2.3.12 warp_Int warp_SetCFactor (warp_Core core, warp_Int level, warp_Int cfactor)

Set the coarsening factor *cfactor* on grid *level* (level 0 is the finest grid). The default factor is 2 on all levels. To change the default factor, use *level* = -1.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>level</i>	<i>level</i> to set coarsening factor on
<i>cfactor</i>	desired coarsening factor

8.2.3.13 warp_Int warp_SetFMG (warp_Core core)

Once called, Warp will use FMG (i.e., F-cycles).

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
-------------	-------------------------------

8.2.3.14 warp_Int warp_SetLooseTol (warp_Core core, warp_Int level, warp_Real loose_tol)

Set loose stopping tolerance *loose_tol* for spatial solves on grid *level* (level 0 is the finest grid).

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>level</i>	<i>level</i> to set <i>loose_tol</i>
<i>loose_tol</i>	tolerance to set

8.2.3.15 warp_Int warp_SetMaxCoarse (warp_Core core, warp_Int max_coarse)

Set max allowed coarse grid size (in terms of C-points)

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>max_coarse</i>	maximum coarse grid size

8.2.3.16 warp_Int warp_SetMaxIter (warp_Core core, warp_Int max_iter)

Set max number of multigrid iterations.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>max_iter</i>	maximum iterations to allow

8.2.3.17 warp_Int warp_SetMaxLevels (warp_Core core, warp_Int max_levels)

Set max number of multigrid levels.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>max_levels</i>	maximum levels to allow

8.2.3.18 warp_Int warp_SetNFMGVcyc (warp_Core *core*, warp_Int *nfmg_Vcyc*)

Set number of V-cycles to use at each FMG level (standard is 1)

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>nfmg_Vcyc</i>	number of V-cycles to do each FMG level

8.2.3.19 warp_Int warp_SetNRelax (warp_Core *core*, warp_Int *level*, warp_Int *nrelax*)

Set the number of relaxation sweeps *nrelax* on grid *level* (level 0 is the finest grid). The default is 1 on all levels. To change the default factor, use *level* = -1. One sweep is a CF relaxation sweep.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>level</i>	<i>level</i> to set <i>nrelax</i> on
<i>nrelax</i>	number of relaxations to do on <i>level</i>

8.2.3.20 warp_Int warp_SetPrintFile (warp_Core *core*, const char * *printfile_name*)

Set output file for runtime print messages. Level of printing is controlled by [warp_SetPrintLevel](#). Default is stdout.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>printfile_name</i>	output file for Warp runtime output

8.2.3.21 warp_Int warp_SetPrintLevel (warp_Core *core*, warp_Int *print_level*)

Set print level for warp. This controls how much information is printed to the Warp print file ([warp_SetPrintFile](#)).

- Level 0: no output
- Level 1: print typical information like a residual history, number of levels in the Warp hierarchy, and so on.
- Level 2: level 1 output, plus debug level output.

Default is level 1.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>print_level</i>	desired print level

8.2.3.22 warp_Int warp_SetRelTol (warp_Core *core*, warp_Real *rtol*)

Set relative stopping tolerance, relative to the initial residual. Be careful. If your initial guess is all zero, then the initial residual may only be nonzero over one or two time values, and this will skew the relative tolerance. Absolute tolerances are recommended.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>rtol</i>	relative stopping tolerance

8.2.3.23 warp_Int warp_SetSpatialCoarsen (warp_Core *core*, warp_PtFcnCoarsen *coarsen*)

Set spatial coarsening routine with user-defined routine. Default is no spatial refinement or coarsening.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>coarsen</i>	function pointer to spatial coarsening routine

8.2.3.24 warp_Int warp_SetSpatialRefine (warp_Core *core*, warp_PtFcnRefine *refine*)

Set spatial refinement routine with user-defined routine. Default is no spatial refinement or coarsening.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>refine</i>	function pointer to spatial refinement routine

8.2.3.25 warp_Int warp_SetTightxTol (warp_Core *core*, warp_Int *level*, warp_Real *tight_tol*)

Set tight stopping tolerance *tight_tol* for spatial solves on grid *level* (level 0 is the finest grid).

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>level</i>	level to set <i>tight_tol</i>
<i>tight_tol</i>	tolerance to set

8.2.3.26 warp_Int warp_SetWriteLevel (warp_Core *core*, warp_Int *write_level*)

Set write level for warp. This controls how often the user's write routine is called.

- Level 0: Never call the user's write routine
- Level 1: Only call the user's write routine after Warp is finished
- Level 2: Call the user's write routine every iteration and on every level. This is during _warp_FRestrict, during the down-cycle part of a Warp iteration.

Default is level 1.

Parameters

<i>core</i>	warp_Core (_warp_Core) struct
<i>write_level</i>	desired write_level

8.2.3.27 warp_Int warp_SplitCommworld (const MPI_Comm * *comm_world*, warp_Int *px*, MPI_Comm * *comm_x*, MPI_Comm * *comm_t*)

Split MPI commworld into *comm_x* and *comm_t*, the spatial and temporal communicators. The total number of processors will equal Px*Pt, where Px is the number of procs in space, and Pt is the number of procs in time.

Parameters

<i>comm_world</i>	Global communicator to split
<i>px</i>	Number of processors parallelizing space for a single time step
<i>comm_x</i>	Spatial communicator (written as output)
<i>comm_t</i>	Temporal communicator (written as output)

8.3 Warp test routines

Functions

- `warp_Int warp_TestInitWrite (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free)`
- `warp_Int warp_TestClone (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone)`
- `warp_Int warp_TestSum (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum)`
- `warp_Int warp_TestDot (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot)`
- `warp_Int warp_TestBuf (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack)`
- `warp_Int warp_TestCoarsenRefine (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`
- `warp_Int warp_TestAll (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`

8.3.1 Detailed Description

These are sanity check routines to help a user test their Warp code.

8.3.2 Function Documentation

8.3.2.1 `warp_Int warp_TestAll (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`

Runs all of the individual warp_Test* routines

- Returns 0 if the tests fail
- Returns 1 if the tests pass
- Check the log messages to see details of which tests failed.

Parameters

<code>app</code>	User defined App structure
<code>comm_x</code>	Spatial communicator
<code>fp</code>	File pointer (could be stdout or stderr) for log messages
<code>t</code>	Time value to initialize test vectors with

<i>fdt</i>	Fine time step value that you spatially coarsen from
<i>cdt</i>	Coarse time step value that you coarsen to
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>free</i>	Free a warp_Vector
<i>clone</i>	Clone a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors
<i>dot</i>	Compute dot product of two warp_Vectors
<i>bufsize</i>	Computes size in bytes for one warp_Vector MPI buffer
<i>bufpack</i>	Packs MPI buffer to contain one warp_Vector
<i>bufunpack</i>	Unpacks MPI buffer into a warp_Vector
<i>coarsen</i>	Spatially coarsen a vector. If NULL, test is skipped.
<i>refine</i>	Spatially refine a vector. If NULL, test is skipped.

8.3.2.2 `warp_Int warp_TestBuf (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack)`

Test the BufPack, BufUnpack and BufSize functions.

A vector is initialized at time *t*, packed into a buffer, then unpacked from a buffer. The unpacked result must equal the original vector.

- Returns 0 if the tests fail
- Returns 1 if the tests pass
- Check the log messages to see details of which tests failed.

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to test Buffer routines (used to initialize the vectors)
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>free</i>	Free a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors
<i>dot</i>	Compute dot product of two warp_Vectors
<i>bufsize</i>	Computes size in bytes for one warp_Vector MPI buffer
<i>bufpack</i>	Packs MPI buffer to contain one warp_Vector
<i>bufunpack</i>	Unpacks MPI buffer containing one warp_Vector

8.3.2.3 `warp_Int warp_TestClone (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone)`

Test the clone function.

A vector is initialized at time *t*, cloned, and both vectors are written. Then both vectors are free-d. The user is to check (via the write function) to see if it is identical.

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to test clone with (used to initialize the vectors)
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>write</i>	Write a warp_Vector (can be NULL for no writing)
<i>free</i>	Free a warp_Vector
<i>clone</i>	Clone a warp_Vector

8.3.2.4 `warp_Int warp_TestCoarsenRefine (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`

Test the Coarsen and Refine functions.

A vector is initialized at time *t*, and various sanity checks on the spatial coarsening and refinement routines are run.

- Returns 0 if the tests fail
- Returns 1 if the tests pass
- Check the log messages to see details of which tests failed.

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to initialize test vectors
<i>fdt</i>	Fine time step value that you spatially coarsen from
<i>cdt</i>	Coarse time step value that you coarsen to
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>write</i>	Write a warp_Vector (can be NULL for no writing)
<i>free</i>	Free a warp_Vector
<i>clone</i>	Clone a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors
<i>dot</i>	Compute dot product of two warp_Vectors
<i>coarsen</i>	Spatially coarsen a vector
<i>refine</i>	Spatially refine a vector

8.3.2.5 `warp_Int warp_TestDot (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot)`

Test the dot function.

A vector is initialized at time *t* and then cloned. Various dot products like $\langle 3 v, v \rangle / \langle v, v \rangle$ are computed with known output, e.g., $\langle 3 v, v \rangle / \langle v, v \rangle$ must equal 3. If all the tests pass, then 1 is returned.

- Returns 0 if the tests fail
- Returns 1 if the tests pass
- Check the log messages to see details of which tests failed.

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to test Dot with (used to initialize the vectors)
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>free</i>	Free a warp_Vector
<i>clone</i>	Clone a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors
<i>dot</i>	Compute dot product of two warp_Vectors

8.3.2.6 `warp_Int warp_TestInitWrite (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free)`

Test the init, write and free functions.

A vector is initialized at time *t*, written, and then free-d

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to test init with (used to initialize the vectors)
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>write</i>	Write a warp_Vector (can be NULL for no writing)
<i>free</i>	Free a warp_Vector

8.3.2.7 `warp_Int warp_TestSum (warp_App app, MPI_Comm comm_x, FILE * fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum)`

Test the sum function.

A vector is initialized at time *t*, cloned, and then these two vectors are summed a few times, with the results written. The vectors are then free-d. The user is to check (via the write function) that the output matches the sum of the two original vectors.

Parameters

<i>app</i>	User defined App structure
<i>comm_x</i>	Spatial communicator
<i>fp</i>	File pointer (could be stdout or stderr) for log messages
<i>t</i>	Time value to test Sum with (used to initialize the vectors)
<i>init</i>	Initialize a warp_Vector on finest temporal grid
<i>write</i>	Write a warp_Vector (can be NULL for no writing)
<i>free</i>	Free a warp_Vector
<i>clone</i>	Clone a warp_Vector
<i>sum</i>	Compute vector sum of two warp_Vectors

9 File Documentation

9.1 warp.h File Reference

Typedefs

- `typedef int warp_Int`
- `typedef double warp_Real`
- `typedef struct _warp_App_struct * warp_App`
- `typedef struct
 _warp_Vector_struct * warp_Vector`
- `typedef struct
 _warp_Status_struct * warp_Status`
- `typedef warp_Int(* warp_PtFcnPhi)(warp_App app, warp_Real tstart, warp_Real tstop, warp_Real accuracy, warp_Vector u, warp_Int *rfactor_ptr)`
- `typedef warp_Int(* warp_PtFcnInit)(warp_App app, warp_Real t, warp_Vector *u_ptr)`
- `typedef warp_Int(* warp_PtFcnClone)(warp_App app, warp_Vector u, warp_Vector *v_ptr)`
- `typedef warp_Int(* warp_PtFcnFree)(warp_App app, warp_Vector u)`
- `typedef warp_Int(* warp_PtFcnSum)(warp_App app, warp_Real alpha, warp_Vector x, warp_Real beta, warp_Vector y)`
- `typedef warp_Int(* warp_PtFcnDot)(warp_App app, warp_Vector u, warp_Vector v, warp_Real *dot_ptr)`
- `typedef warp_Int(* warp_PtFcnWrite)(warp_App app, warp_Real t, warp_Status status, warp_Vector u)`
- `typedef warp_Int(* warp_PtFcnBufSize)(warp_App app, warp_Int *size_ptr)`
- `typedef warp_Int(* warp_PtFcnBufPack)(warp_App app, warp_Vector u, void *buffer)`
- `typedef warp_Int(* warp_PtFcnBufUnpack)(warp_App app, void *buffer, warp_Vector *u_ptr)`
- `typedef warp_Int(* warp_PtFcnCoarsen)(warp_App app, warp_Real tstart, warp_Real f_tminus, warp_Real f_tplus, warp_Real c_tminus, warp_Real c_tplus, warp_Vector fu, warp_Vector *cu_ptr)`
- `typedef warp_Int(* warp_PtFcnRefine)(warp_App app, warp_Real tstart, warp_Real f_tminus, warp_Real f_tplus, warp_Real c_tminus, warp_Real c_tplus, warp_Vector cu, warp_Vector *fu_ptr)`
- `typedef struct _warp_Core_struct * warp_Core`

Functions

- `warp_Int warp_Init (MPI_Comm comm_world, MPI_Comm comm, warp_Real tstart, warp_Real tstop, warp_Int ntime, warp_App app, warp_PtFcnPhi phi, warp_PtFcnInit init, warp_PtFcnClone clone, warp_PtFcnFree free, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnWrite write, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack, warp_Core *core_ptr)`
- `warp_Int warp_Drive (warp_Core core)`
- `warp_Int warp_Destroy (warp_Core core)`
- `warp_Int warp_PrintStats (warp_Core core)`
- `warp_Int warp_SetLoosexTol (warp_Core core, warp_Int level, warp_Real loose_tol)`
- `warp_Int warp_SetTightxTol (warp_Core core, warp_Int level, warp_Real tight_tol)`
- `warp_Int warp_SetMaxLevels (warp_Core core, warp_Int max_levels)`
- `warp_Int warp_SetMaxCoarse (warp_Core core, warp_Int max_coarse)`
- `warp_Int warp_SetAbsTol (warp_Core core, warp_Real atol)`
- `warp_Int warp_SetRelTol (warp_Core core, warp_Real rtol)`
- `warp_Int warp_SetNRelax (warp_Core core, warp_Int level, warp_Int nrelax)`
- `warp_Int warp_SetCFactor (warp_Core core, warp_Int level, warp_Int cfactor)`
- `warp_Int warp_SetMaxIter (warp_Core core, warp_Int max_iter)`
- `warp_Int warp_SetFMG (warp_Core core)`
- `warp_Int warp_SetNFMGVcyc (warp_Core core, warp_Int nfmvg_Vcyc)`

- `warp_Int warp_SetSpatialCoarsen (warp_Core core, warp_PtFcnCoarsen coarsen)`
- `warp_Int warp_SetSpatialRefine (warp_Core core, warp_PtFcnRefine refine)`
- `warp_Int warp_SetPrintLevel (warp_Core core, warp_Int print_level)`
- `warp_Int warp_SetPrintFile (warp_Core core, const char *printfile_name)`
- `warp_Int warp_SetWriteLevel (warp_Core core, warp_Int write_level)`
- `warp_Int warp_SplitCommworld (const MPI_Comm *comm_world, warp_Int px, MPI_Comm *comm_x, MPI_Comm *comm_t)`
- `warp_Int warp_GetStatusResidual (warp_Status status, warp_Real *rnorm_ptr)`
- `warp_Int warp_GetStatusIter (warp_Status status, warp_Int *iter_ptr)`
- `warp_Int warp_GetStatusLevel (warp_Status status, warp_Int *level_ptr)`
- `warp_Int warp_GetStatusDone (warp_Status status, warp_Int *done_ptr)`
- `warp_Int warp_GetNumIter (warp_Core core, warp_Int *niter_ptr)`
- `warp_Int warp_GetRNorm (warp_Core core, warp_Real *rnorm_ptr)`

9.1.1 Detailed Description

Define headers for user interface routines. This file contains routines used to allow the user to initialize, run and get and set warp.

9.1.2 Typedef Documentation

9.1.2.1 `typedef int warp_Int`

Defines integer type

9.1.2.2 `typedef double warp_Real`

Defines floating point type

9.2 warp_test.h File Reference

Functions

- `warp_Int warp_TestInitWrite (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free)`
- `warp_Int warp_TestClone (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone)`
- `warp_Int warp_TestSum (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum)`
- `warp_Int warp_TestDot (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot)`
- `warp_Int warp_TestBuf (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack)`
- `warp_Int warp_TestCoarsenRefine (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnWrite write, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`
- `warp_Int warp_TestAll (warp_App app, MPI_Comm comm_x, FILE *fp, warp_Real t, warp_Real fdt, warp_Real cdt, warp_PtFcnInit init, warp_PtFcnFree free, warp_PtFcnClone clone, warp_PtFcnSum sum, warp_PtFcnDot dot, warp_PtFcnBufSize bufsize, warp_PtFcnBufPack bufpack, warp_PtFcnBufUnpack bufunpack, warp_PtFcnCoarsen coarsen, warp_PtFcnRefine refine)`

9.2.1 Detailed Description

Define headers for Warp test routines. This file contains routines used to test a user's Warp wrapper routines one-by-one.

Index

User interface routines, 19
 warp_Core, 19
 warp_Destroy, 20
 warp_Drive, 21
 warp_GetNumIter, 21
 warp_GetRNorm, 21
 warp_GetStatusDone, 21
 warp_GetStatusIter, 21
 warp_GetStatusLevel, 21
 warp_GetStatusResidual, 22
 warp_Init, 22
 warp_PrintStats, 22
 warp_SetAbsTol, 22
 warp_SetCFactor, 24
 warp_SetFMG, 24
 warp_SetLooseTol, 24
 warp_SetMaxCoarse, 24
 warp_SetMaxIter, 24
 warp_SetMaxLevels, 24
 warp_SetNFMGVcyc, 25
 warp_SetNRelax, 25
 warp_SetPrintFile, 25
 warp_SetPrintLevel, 25
 warp_SetRelTol, 25
 warp_SetSpatialCoarsen, 26
 warp_SetSpatialRefine, 26
 warp_SetTightxTol, 26
 warp_SetWriteLevel, 26
 warp_SplitCommworld, 26

User-written routines, 16
 warp_App, 17
 warp_PtFcnBufPack, 17
 warp_PtFcnBufSize, 17
 warp_PtFcnBufUnpack, 17
 warp_PtFcnClone, 17
 warp_PtFcnCoarsen, 17
 warp_PtFcnDot, 18
 warp_PtFcnFree, 18
 warp_PtFcnInit, 18
 warp_PtFcnPhi, 18
 warp_PtFcnRefine, 18
 warp_PtFcnSum, 18
 warp_PtFcnWrite, 18
 warp_Status, 18
 warp_Vector, 18

Warp test routines, 28
 warp_TestAll, 28
 warp_TestBuf, 29
 warp_TestClone, 29
 warp_TestCoarsenRefine, 30
 warp_TestDot, 30
 warp_TestInitWrite, 31
 warp_TestSum, 31
 warp.h, 32
 warp_Int, 33
 warp_Real, 33
 warp_App
 User-written routines, 17
 warp_Core
 User interface routines, 19
 warp_Destroy
 User interface routines, 20
 warp_Drive
 User interface routines, 21
 warp_GetNumIter
 User interface routines, 21
 warp_GetRNorm
 User interface routines, 21
 warp_GetStatusDone
 User interface routines, 21
 warp_GetStatusIter
 User interface routines, 21
 warp_GetStatusLevel
 User interface routines, 21
 warp_Init
 User interface routines, 22
 warp_Int
 warp.h, 33
 warp_PrintStats
 User interface routines, 22
 warp_PtFcnBufPack
 User-written routines, 17
 warp_PtFcnBufSize
 User-written routines, 17
 warp_PtFcnBufUnpack
 User-written routines, 17
 warp_PtFcnClone
 User-written routines, 17
 warp_PtFcnCoarsen
 User-written routines, 17
 warp_PtFcnDot
 User-written routines, 18
 warp_PtFcnFree
 User-written routines, 18
 warp_PtFcnInit
 User-written routines, 18
 warp_PtFcnPhi
 User-written routines, 18
 warp_PtFcnRefine

User-written routines, 18
warp_PtFcnSum
 User-written routines, 18
warp_PtFcnWrite
 User-written routines, 18
warp_Real
 warp.h, 33
warp_SetAbsTol
 User interface routines, 22
warp_SetCFactor
 User interface routines, 24
warp_SetFMG
 User interface routines, 24
warp_SetLooseTol
 User interface routines, 24
warp_SetMaxCoarse
 User interface routines, 24
warp_SetMaxIter
 User interface routines, 24
warp_SetMaxLevels
 User interface routines, 24
warp_SetNFMGVcyc
 User interface routines, 25
warp_SetNRelax
 User interface routines, 25
warp_SetPrintFile
 User interface routines, 25
warp_SetPrintLevel
 User interface routines, 25
warp_SetRelTol
 User interface routines, 25
warp_SetSpatialCoarsen
 User interface routines, 26
warp_SetSpatialRefine
 User interface routines, 26
warp_SetTightxTol
 User interface routines, 26
warp_SetWriteLevel
 User interface routines, 26
warp_SplitCommworld
 User interface routines, 26
warp_Status
 User-written routines, 18
warp_TestAll
 Warp test routines, 28
warp_TestBuf
 Warp test routines, 29
warp_TestClone
 Warp test routines, 29
warp_TestCoarsenRefine
 Warp test routines, 30
warp_TestDot
 Warp test routines, 30
warp_TestInitWrite

 Warp test routines, 31
warp_TestSum
 Warp test routines, 31
warp_Vector
 User-written routines, 18
warp_test.h, 33